

# Assignment 3

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

import tensorflow as tf

from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense

from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

import cv2
```

```
In [ ]: data_path = './flowers/'
batch_size = 32
target_size = (64, 64)
```

```
In [ ]: train_datagen = ImageDataGenerator(rescale=1./255,
                                           shear_range=0.2,
                                           zoom_range=0.2,
                                           width_shift_range=0.1,
                                           height_shift_range=0.1,
                                           horizontal_flip=True,
                                           validation_split=0.2)

test_datagen = ImageDataGenerator(rescale=1. / 255, validation_split=0.2)
```

```
In [ ]: X_train = train_datagen.flow_from_directory(data_path,
                                                    target_size=target_size,
                                                    batch_size=batch_size,
                                                    subset="training",
                                                    class_mode='categorical')

X_test = test_datagen.flow_from_directory(data_path,
                                           target_size=target_size,
                                           batch_size=batch_size,
                                           subset="validation",
                                           class_mode='categorical')
```

Found 3457 images belonging to 5 classes.  
Found 860 images belonging to 5 classes.

```
In [ ]: model = Sequential()
```

```
In [ ]: model.add(Convolution2D(32, (3, 3), input_shape=(64, 64, 3), activation=
```

```

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(units=64, activation='relu'))
model.add(Dense(units=5, activation='softmax'))

```

In [ ]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147520
dense_1 (Dense)	(None, 5)	325

=====  
 Total params: 176,485  
 Trainable params: 176,485  
 Non-trainable params: 0  
 =====

In [ ]: `model.compile(optimizer='adam', loss='categorical_crossentropy', metrics:`

In [ ]: `model.fit(X_train, steps_per_epoch=100, epochs=20)`

```

Epoch 1/20
100/100 [=====] - 8s 74ms/step - loss: 0.8592 - accuracy: 0.6693
Epoch 2/20
100/100 [=====] - 8s 75ms/step - loss: 0.8127 - accuracy: 0.6841
Epoch 3/20
100/100 [=====] - 7s 75ms/step - loss: 0.7724 - accuracy: 0.7009
Epoch 4/20

```

```

100/100 [=====] - 7s 73ms/step - loss: 0.7421 -
accuracy: 0.7198
Epoch 5/20
100/100 [=====] - 8s 78ms/step - loss: 0.7266 -
accuracy: 0.7204
Epoch 6/20
100/100 [=====] - 8s 75ms/step - loss: 0.7393 -
accuracy: 0.7160
Epoch 7/20
100/100 [=====] - 8s 82ms/step - loss: 0.7395 -
accuracy: 0.7163
Epoch 8/20
100/100 [=====] - 8s 80ms/step - loss: 0.6993 -
accuracy: 0.7251
Epoch 9/20
100/100 [=====] - 8s 76ms/step - loss: 0.6899 -
accuracy: 0.7444
Epoch 10/20
100/100 [=====] - 8s 76ms/step - loss: 0.6781 -
accuracy: 0.7457
Epoch 11/20
100/100 [=====] - 8s 76ms/step - loss: 0.6612 -
accuracy: 0.7463
Epoch 12/20
100/100 [=====] - 8s 76ms/step - loss: 0.6526 -
accuracy: 0.7539
Epoch 13/20
100/100 [=====] - 8s 76ms/step - loss: 0.6387 -
accuracy: 0.7485
Epoch 14/20
100/100 [=====] - 8s 77ms/step - loss: 0.6011 -
accuracy: 0.7712
Epoch 15/20
100/100 [=====] - 8s 76ms/step - loss: 0.5926 -
accuracy: 0.7826
Epoch 16/20
100/100 [=====] - 8s 80ms/step - loss: 0.6164 -
accuracy: 0.7608
Epoch 17/20
100/100 [=====] - 8s 77ms/step - loss: 0.6137 -
accuracy: 0.7690
Epoch 18/20
100/100 [=====] - 8s 75ms/step - loss: 0.5976 -
accuracy: 0.7703
Epoch 19/20
100/100 [=====] - 8s 80ms/step - loss: 0.6437 -
accuracy: 0.7545
Epoch 20/20
100/100 [=====] - 8s 76ms/step - loss: 0.6004 -
accuracy: 0.7700

```

Out[ ]: <keras.callbacks.History at 0x274702feb90>

```
In [ ]: model.save("model.h5")
```

```
In [ ]: def predict():
        img = image.load_img("./rose.jpg", target_size=target_size)
        x = image.img_to_array(img)
        x = tf.expand_dims(x, 0)

        labels = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']
```

```
pred = model.predict(x)
prediction = labels[np.argmax(pred[0])]

print(f'The given image is a {prediction}')
plt.imshow(plt.imread("./rose.jpg"))
plt.axis('off')
plt.show()
```

In [ ]:

```
predict()
```

1/1 [=====] - 0s 22ms/step

The given image is a rose

